

Intro to Logistic Regression

Kevin Cummiskey

November 18, 2019

```
library(tidyverse)
```

Section 6.2 Introduction to Logistic Regression

Review

In the article “The incidence of thyroid disorders in the community: A twenty-year follow-up of the Wickham survey” by Vanderpump et al., 443 of the 582 smokers and 502 of the 732 nonsmokers were still alive at the 20 year follow-up.

Calculate the odds ratio and logs odds ratio of being alive comparing smokers to nonsmokers.

Perform the chi-square test to determine if there is a significant association between smoking and survival. State the appropriate hypotheses and report the chi-square test statistic and p -value.

Does your result above mean smoking raises the probability of being alive?

Useful log rules for this lesson

- $e^{\ln r} = r$
- $\ln(e^r) = r$
- $e^r \times e^s = e^{r+s}$
- $\ln\left(\frac{r}{s}\right) = \ln(r) - \ln(s)$

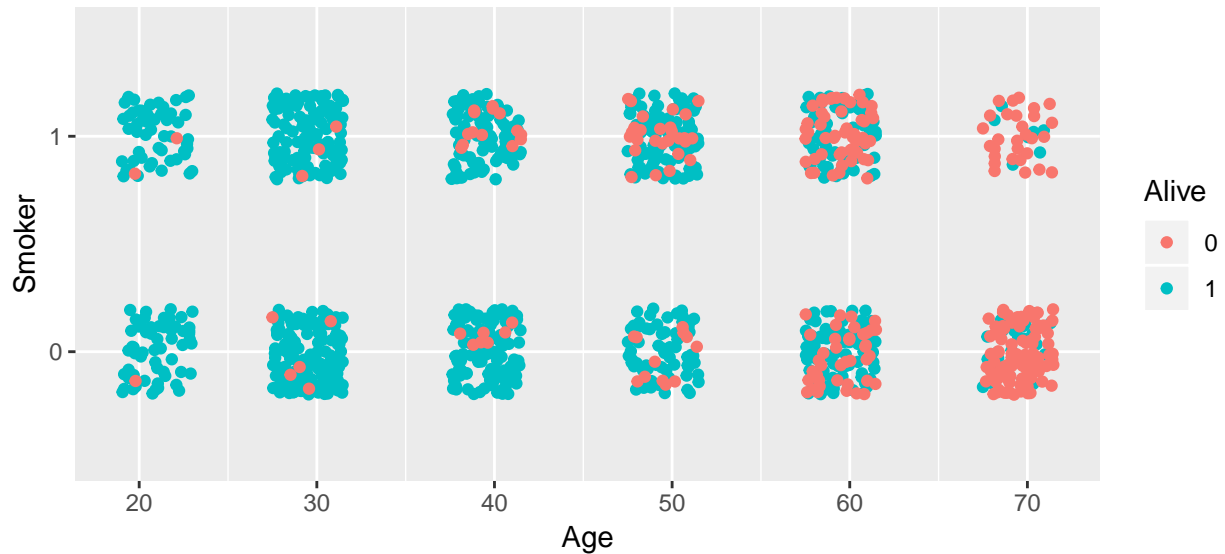
Introduction to Logistic Regression

```
smoke = read.table(file = "smoke.csv", header = T, sep = ",")
smoke$Alive = factor(smoke$Alive)
smoke$Smoker = factor(smoke$Smoker)
smoke %>% group_by(Smoker,Alive) %>% count() %>% spread(key = "Smoker",value = "n")
```

```
## # A tibble: 2 x 3
## # Groups:   Alive [2]
##   Alive   `0`   `1`
##   <fct> <int> <int>
## 1 0      166   126
```

```
## 2 1      502   443
```

```
smoke %>% ggplot(aes(x = Age, y = Smoker, color = Alive)) + geom_jitter(width = 2, height = 0.2)
```




Based on the plot, is age a confounding variable of the association between smoking and survival? Explain.

Let's look at the relationship between survival and age.

```
#Let's look at a plot of age vs proportion alive
summary = smoke %>% count(Age, Alive) %>%
  group_by(Age) %>%
  mutate(Alive_proportion = n/sum(n),
         total = sum(n)) %>%
  filter(Alive == 1) %>% select(-c(Alive,n))


summary %>%
  ggplot(aes(x = Age, y = Alive_proportion)) +
  geom_point()
```



Explain two reasons a linear model for age and alive proportion is not appropriate?

Let's try a log odds (or *logit*) transformation.

```
summary = summary %>%
  mutate(logOdds = log(Alive_proportion/(1-Alive_proportion)))
summary %>% ggplot(aes(x = Age, y = logOdds)) +
  geom_point()
```



Is a linear model on the log odds appropriate? How does the logit transformation ensure probabilities are between 0 and 1?

Let's fit the following model.

$$\log\left(\frac{\pi_i}{1 - \pi_i}\right) = \beta_0 + \beta_1 \text{Age}_i$$

How do we interpret β_0 and β_1 ? Would we expect β_1 to be positive or negative? Why is there no ϵ_i on this model?

```
# In practice, we would just do this:
model_age = glm(Alive ~ Age, data = smoke, family = "binomial")
summary(model_age)

##
## Call:
## glm(formula = Alive ~ Age, family = "binomial", data = smoke)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.0327   0.1422   0.2295   0.6782   1.5837
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  6.972626   0.410228   17.0    <2e-16 ***
## Age         -0.113535   0.007232  -15.7    <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1352.0  on 1236  degrees of freedom
## Residual deviance:  938.4  on 1235  degrees of freedom
## AIC: 942.4
##
## Number of Fisher Scoring iterations: 5
```

What is the predicted odds ratio associated with a one-year increase in age?

What is the predicted odds ratio associated with a ten-year increase in age?

Key idea: the slope of a logistic regression model indicates a multiplicative change in the odds.

What would we conclude from these results?

How do we find predicted probabilities of survival for each subject?

```
summary_model = model_age %>% fortify()
summary_model = summary_model %>% mutate(predicted_prob = exp(.fitted)/(1 + exp(.fitted)))
```

Let's look at the relationship between smoking and survival

```
model_smoker = glm(Alive ~ Smoker, data = smoke, family = "binomial")
summary(model_smoker)
```

```
##
## Call:
## glm(formula = Alive ~ Smoker, family = "binomial", data = smoke)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.7364   0.7076   0.7076   0.7559   0.7559
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  1.10661    0.08953  12.360  <2e-16 ***
## Smoker1      0.15068    0.13494   1.117   0.264
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1352.0  on 1236  degrees of freedom
## Residual deviance: 1350.8  on 1235  degrees of freedom
## AIC: 1354.8
##
## Number of Fisher Scoring iterations: 4
```

How do we interpret β_0 and β_1 ? Have we seen these values before?